

Minimizing Faulty Executions of Distributed Systems

Colin Scott, Aurojit Panda, Vjekoslav Brajkovic,
George Necula, Arvind Krishnamurthy, Scott Shenker

Motivation in <60s

- ▶ Distributed systems are highly complex (concurrency, asynchrony, partial failure, ...)
- ▶ \therefore Distributed systems are bug prone
- ▶ Popular way to catch bugs: QA (fuzz) testing
- ▶ Failing QA tests often hard to understand (hours long, multiple GBs in size)
- ▶ Much less time consuming to debug small traces

Dist'Sys Computational Model*

Distributed System: Collection of N processes

Each process p :

- ▶ Has unbounded internal storage
- ▶ Starts in a known initial state
- ▶ Changes states deterministically according to a transition function

*Inspired by FLP, JACM '85

Dist'Sys Computational Model

A message is a pair: (p, m)

The network maintains a buffer of sent but not yet delivered messages

Dist'Sys Computational Model

A configuration consists of:

- ▶ The internal state of each process
- ▶ The contents of the network buffer

Initially: empty network buffer, known initial states

Dist'Sys Computational Model

A *step* moves the system from one config to another

Normally:

- ▶ Network delivers some pending message (p, m)
- ▶ p enters a new state according to old state & transition function
- ▶ p sends a finite set of messages to other processes*

*May include *timer* messages to be delivered to itself later

Dist'Sys Computational Model

A *step* can also take another form: *external events*

Either:

- ▶ External message is delivered
- ▶ Process p fails, i.e. has its internal state reinitialized

Note: no need to explicitly model network partitions

Dist'Sys Computational Model

A *schedule* τ is a finite sequence of events (either internal messages or external events) that can be applied in turn starting from the initial configuration.

An *execution* is the corresponding sequence of steps.

Faulty executions

An invariant is a predicate P (a safety condition) over the internal state of each process at a particular configuration C .

A faulty execution is one that ends in $P(C) = \textit{false}$

Problem Statement

Given: P and a schedule τ that produces an execution that violates P

Primary goal: find schedule τ' s.t.

- ▶ $|\tau'| \leq |\tau|$
- ▶ if we remove any single external event e from τ' ,
- ▶ $\neg \exists \tau''$ containing same external events - e , s.t.
 τ'' violates P

Problem Statement

Secondary goal:

After finding τ' , remove as many internal message events from τ' as possible while still causing violation

General Approach

Repeat:

1. Pick subsequence of external events from τ
2. Try to find a fault-inducing schedule containing those external events
3. If found: recurse

Key Challenge

Step 2 is intractable!

Research Agenda

- ▶ Find *heuristics* for exploring the schedule space s.t. fault-inducing schedules are found quickly
- ▶ Develop heuristics by *experimenting* with faulty executions found in real systems
- ▶ Formally characterize what program properties hold in those systems s.t. heuristics are effective

Contributions vs. previous paper

- ▶ Cleaner, more general problem statement
- ▶ New tool gives better control over the execution
- ▶ Broader system target (Spark, Raft, Pastry, JS)
- ▶ Novel heuristics, empirically validated
- ▶ Conjectures for explanatory program properties

New Tool

- ▶ Interpose on the Akka actor system
- ▶ Actors:
 - ▶ are single threaded, no shared state
 - ▶ compute by: i. receiving a message, ii. changing state, iii. sending out messages. (Atomic steps!)
 - ▶ send timer messages rather than reading from clock. (No dependence on wall-clock time!)

New Tool

- ▶ The Akka runtime is our “network”
- ▶ Interposition allows us to completely specify a linearizable sequence of steps to take
- ▶ Start with a test scenario written on top of Akka, then systematically explore schedules for that scenario

System Target

- ▶ akka-raft: consensus protocol written on top of akka
- ▶ Spark: control plane built on top of Akka
- ▶ Pastry: DHT written on top of kaka
- ▶ JavaScript: separate tool, similar computational model

System Target

- ▶ Find a variety of bug types in these systems
(holy grail: find heuristics that are applicable to a wide range of bugs)
- ▶ If possible, reproduce known bugs
- ▶ Otherwise, find new bugs via fuzz testing

Example Bug

- ▶ akka-raft:
 - ▶ Two leaders elected in the same term
 - ▶ <Whiteboard>

Previous Heuristics

- ▶ Pick one schedule per subsequence
- ▶ Define functional equivalence (fingerprints)
- ▶ If unexpected event during replay, ignore
- ▶ If absent event during replay, skip over

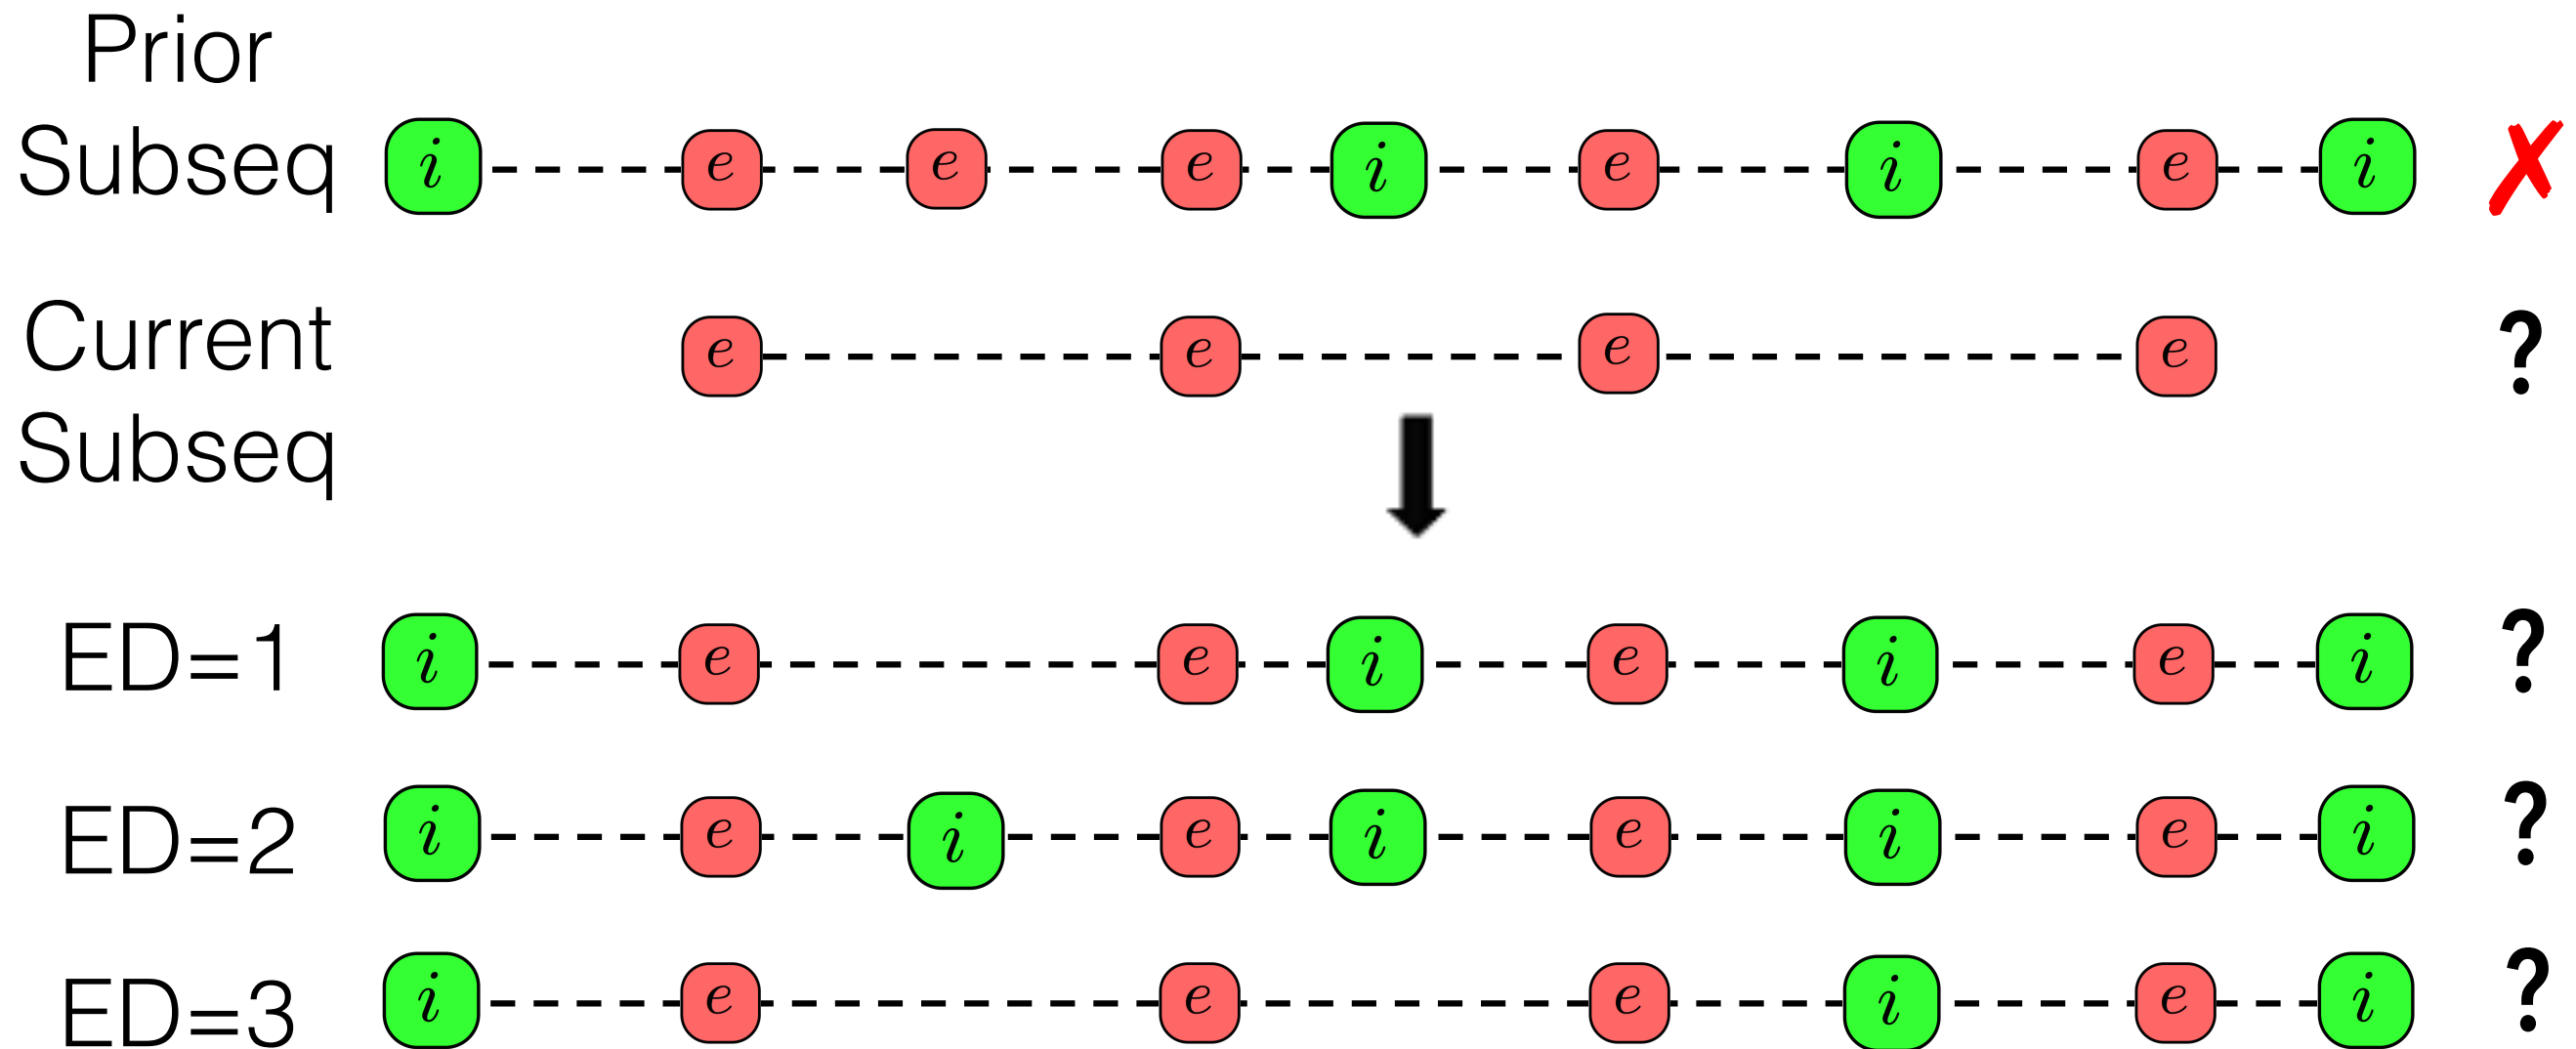
Results w/ Previous Heuristics

	Short Run	Long Run
Original Messages Deliveries	90	360
Removed by Provenance	26	25
Removed by DDmin	9	62
Removed by internal minimization	20	192
Final # of events	35 (7 external)	81 (8 external)

Old heuristics are suboptimal!

	Short Run	Long Run
Original Messages Deliveries	90	360
Removed by Provenance	26	25
Removed by DDmin	9	62
Removed by internal minimization	20	192
Final # of events	35 (7 external)	81 (8 external)

Heuristic: Bounded Edit Distance



Evaluation Methodology

- ▶ Empirically compare effectiveness of heuristics
- ▶ Key metrics:
 - ▶ Closeness to minimal result (if known)
 - ▶ # of schedules explored before terminating

Program Properties

- ▶ Data Independence
- ▶ Bounded Edit-Distance
- ▶ Commutativity
- ▶ Recency of State

Stop Gap: Fairness

- ▶ Enforce fixed computational budget for every external event subsequence, spread evenly over external events

Related Work

- ▶ traditional minimization (DDmin, QuickCheck's shrinking)
 - ▶ not immediately applicable to systems with intermediate external events
- ▶ best-effort minimization (QuickCheck applied to concurrent systems, field failures)
 - ▶ don't systematically handle concurrency
- ▶ model checking minimization: MAX-Sat, interpolation (ConcBugAssist, DSPs)
 - ▶ many disadvantages of model checking. See Q & A.
- ▶ deterministic replay log minimization (schedule minimization, program analysis)
 - ▶ don't allow divergence during replay! i.e. minimization specification is overly specific -> minimization results aren't great
 - ▶ program flow analysis work is tied to a single language, + high perf overhead
- ▶ model inference, log summarization (synoptic, csight, invariant mining)
 - ▶ summarize the events that occurred, but don't actually minimize the test case. Apply this after you minimize the test case with our techniques.
- ▶ program slicing, automated debugging (experimental, program analysis)
 - ▶ minimize program statements, but not the test case itself. Useful for debugging, but not troubleshooting. Apply this after you minimize the test case with our techniques.

Q & A

- ▶ Why not analyze the program?
 - ▶ ties you to one language
 - ▶ parsers usually aren't complete; will sometimes have to make assumptions that aren't realistic
 - ▶ requires engineering effort to get working
 - ▶ high computational overhead of tracing
 - ▶ by solving a harder problem, we come up with more interesting solutions
- ▶ Is QA testing going to be around forever? Why not use newer more sophisticated testing tools?
 - ▶ As long as you have bugs in production, you will need minimization.